

Bold black for commands/keywords

Blue: explanations

Red: data/table/column names

SQL Course Notes – Updated VER

◆ DDL – Data Definition Language

Commands:

CREATE, ALTER, DROP, RENAME, TRUNCATE

◆ CREATE Statement

```
create table staff (
    civilid number(3),
    staffname varchar(30),
    email varchar(40));
describe staff; → show the structure of the table.
```

◆ ALTER Statement

Add column:

```
alter table student
add gender varchar(6); --> to add new column.
```

-To add a primary key:

```
alter table student
add primary key(id);
```

-Alter Column DataType:

```
ALTER TABLE table_name
MODIFY COLUMN column_name datatype;
```

EX1: **Alter teacher Teacher_name column Drop DOB;**

EX2: **Alter Table Teacher Modify Teacher_name char;**

-Drop column:

```
ALTER TABLE teacher
```

```
DROP COLUMN DOB;
```

-Modify column type:

```
ALTER TABLE teacher
```

```
MODIFY teacher_name CHAR;
```

◆ DML – Data Manipulation Language

Commands:

INSERT, DELETE, UPDATE, SELECT

◆ INSERT Statement

```
INSERT INTO student (id, name, location, mob, gender)
```

```
VALUES (200, 'Ahmed', 'Muscat', 90989876, 'Male');
```

Use single quotes for string (varchar) values.

◆ SELECT Statement

Selects columns or all data from a table.

```
select columnName from Tablename; --> to show the data
```

```
select Id from student;
```

```
select id,gender from student;
```

```
select * from student; --> the * means all
```

String concatenation:

select empid || empname from employee; --> view data of 2 columns together

select * from employee where empid=10; --> will get one record

select * from employee where empname='ahmed'; --> single quotes for strings.

select * from employee where empid > 10; --> all records greater than 10

select * from employee where empid >=10; -->10 and more

select * from teacher where dob = '01-jan-2020'

With alias:

select teacher_name as staff from teacher

◆ Constraints

Types:

- NOT NULL
 - UNIQUE
 - PRIMARY KEY
 - FOREIGN KEY
 - CHECK
-

◆ Column-Level Constraint Example

```
CREATE TABLE empTLC (
    Empno NUMBER(4) CONSTRAINT emp_empno_pk PRIMARY KEY,
    Ename VARCHAR(20) CONSTRAINT emp_ename_nnTL NOT NULL,
    birthday DATE CONSTRAINT emp_birthday_nnTL NOT NULL,
    Salary NUMBER(6) CONSTRAINT emp_salary_nnTL NOT NULL,
    Email VARCHAR(25),
    CONSTRAINT emp_email_ukTL UNIQUE
);
```

◆ Add a CHECK Constraint

```
ALTER TABLE emp1  
MODIFY salary NUMBER(6) CONSTRAINT emp_salary_n33  
CHECK(salary > 2500);
```

◆ SELECT Constraints

```
SELECT * FROM user_constraints;
```

◆ Foreign Key – Column Level

```
create table Emp(  
    Empno    Number(4) constraint emp_empno_pk Primary Key,  
    Ename    Varchar(15) constraint emp_ename_nn Not Null,  
    BirthDate Date constraint emp_bdate_nn Not Null,  
    Salary    Number(6) constraint emp_salary_nn Not Null,  
    Deptno   Number(4) References Dept(deptno)  
); # we add the FK constrain whitin the column it self.
```

◆ Foreign Key – Table Level

```
CREATE TABLE Emp (  
    Deptno NUMBER(4),  
    CONSTRAINT emp_deptno_fk FOREIGN KEY(deptno) REFERENCES Dept(deptno)  
); # we added the constraint at the end.
```

◆ Multiple CHECK Constraints

```
CREATE TABLE personSal (  
    Person_ID INT,  
    Last_Name VARCHAR(255),  
    FirstName VARCHAR(255),
```

```
salary NUMBER(8,2),  
CONSTRAINT Personal_min_sal_ck CHECK (salary > 0),  
CONSTRAINT Personal_max_sal_ck CHECK (salary < 1000)  
);
```

* we can add multiple constraint to the same columns, each with a unique constraint name.

◆ Managing Constraints

- Add constraint:

ALTER table **personSal**

Add constraint **personal_pid_pk** primary key (**person_ID**);

- Drop constraint:

ALTER TABLE **personSal**

DROP CONSTRAINT **personal_max_sal_ck**;

- Disable constraint:

ALTER TABLE **personal**

DISABLE CONSTRAINT **personal_pid_pk**;

- Enable:

ALTER TABLE **employees**

ENABLE CONSTRAINT **emp_id**;

- Modify NOT NULL:

ALTER TABLE **stud**

MODIFY (**email** **CONSTRAINT** **stud_email_nn** **NOT NULL**);

- ◆ **Insert, Update, Delete**

- **Insert record:**

```
INSERT INTO personal1 (person_id, last_name, firstname, salary)  
VALUES (222, 'Ahmed', 'Mohammed', 2000);
```

- **Insert without column list (ensure order matches):**

```
INSERT INTO personal1
```

```
VALUES (333, 'Ali', 'Mohammed', 2000);
```

** order of the values should match columns order in our table.

- **Copy table:**

```
CREATE TABLE copy_person AS
```

```
(SELECT * FROM personal1);
```

- **Update existing data/record:**

```
UPDATE personal1  
SET Salary = 2000  
where person_id= 111;
```

- **Delete value:**

```
DELETE FROM personal1
```

```
WHERE person_id = 111;
```

Deleting multiple values with a condition:

```
DELETE FROM personal1
```

```
WHERE person_id > 250;
```

◆ Boolean Values in SQL (Logic Check Notes)

Can be represented as:

- TRUE / FALSE
 - ON / OFF
 - 1 / 0
 - YES / NO
-

SQL Logical Conditions & Filtering (Where Clause + Operators)

◆ DEFAULT Values

DEFAULT SYSDATE

- Used to assign the current system date to a column, usually in a DATE column.
 - Can be applied when creating or later altering a table.
 - If value is not specified during insert, the default will be used.
-

◆ WHERE Clause – Filtering Rows

Used with:

- Comparison operators
 - Logical operators
 - Pattern matching
 - NULL checks
-

◆ Comparison Operators

- Not Equal

```
SELECT empid, empname, deptid
```

```
FROM employee
```

```
WHERE deptid != 90;
```

- **Equals**

```
SELECT first_name, last_name  
FROM employees  
WHERE last_name = 'Taylor'; ** Case-sensitive search for string values.
```

- **Greater Than or Equal**

```
SELECT first_name, last_name  
FROM my_employee  
WHERE hire_date >= '01-Jan-2000';
```

- **Less Than or Equal**

```
SELECT empid, empname  
FROM employee1  
WHERE salary <= 3000;
```

- **Between**

```
SELECT empid, empname  
FROM employee1  
WHERE salary BETWEEN 500 AND 2500;  
** Avoid rewriting it as salary >= 9000 AND salary <= 11000.
```

- ◊ **IN Operator**

- **Example:**

```
SELECT city, state_province, country_id  
FROM locations  
WHERE country_id IN ('UK', 'CA');
```

◊ LIKE Operator (Pattern Matching)

Select last_name from employee

Where `last_name like 'A%'`; **search for all names start with of A.

Select last_name from employee

Where `last_name like '%d'`; **search for all names Ended with of D

Select last_name from employee

Where `last_name like '%d%`; **search for all names have letter D anywhere in middle.

Select last_name from employee

Where `last_name like '_a%'`; **search for all names have letter A in the second position

Select last_name from employee

Where `last_name like '__a%'`; **search for all names have letter A in the 3rd position

Select last_name from employee

Where `last_name like '____'`; **search for all names have 4 letters

◊ NULL Handling

- IS NOT NULL

`SELECT last_name, manager_id`

`FROM employees`

WHERE `manager_id IS NOT NULL;`

- **IS NULL**

```
SELECT last_name, commission_pct  
FROM employees  
WHERE commission_pct IS NULL;
```

- ◊ **Logical Operators: AND, OR, NOT**

- ◊ **OR Operator**

```
SELECT empid, empname  
FROM employee1  
WHERE empid >= 9000 OR salary <= 11000;
```

*** Returns rows if either condition is true.*

Ex2:

```
SELECT last_name, hire_date, job_id  
FROM employees  
WHERE hire_date > '01-Jan-1998' OR job_id LIKE 'SA%';
```

- ◊ **AND Operator**

```
SELECT empid, empname  
FROM employee1  
WHERE empid >= 9000 AND salary <= 11000;
```

*** Both conditions must be true.*

◊ NOT Operator

```
SELECT department_name, location_id  
FROM employees  
WHERE location_id NOT IN (1700, 1800);
```

```
SELECT *  
FROM personal1  
WHERE NOT last_name = 'Ali';
```

◊ Operator Precedence (AND > OR)

Always evaluated as:

** AND → before OR

```
SELECT *  
FROM personal1  
WHERE last_name = 'Ahmed' AND person_id = 222 OR salary = 333;
```

```
SELECT *  
FROM personal1  
WHERE last_name = 'Ahmed' OR person_id = 227 AND salary = 333;  
** To avoid confusion, use parentheses for clarity.
```

Ordering, Aliasing, and SQL Single Row Functions

◆ Ordering Data

- Ascending Order (default):

```
SELECT last_name, hire_date
```

```
FROM employees
```

```
ORDER BY last_name;
```

- Descending Order:

```
SELECT last_name, hire_date
```

```
FROM employees
```

```
ORDER BY last_name DESC;
```

With Column Aliases:

```
SELECT last_name AS "Name", hire_date AS "Date Started"
```

```
FROM employees
```

```
ORDER BY "Name", hire_date;
```

- Multiple Columns:

```
SELECT department_id, last_name
```

```
FROM employees
```

```
WHERE department_id <= 50
```

```
ORDER BY department_id, last_name DESC;
```

CH6: SQL Functions

◆ DUAL Table

```
SELECT (2 + 2) FROM dual;
```

*** The DUAL table is a dummy table used for calculations or function testing.*

◆ Character Functions

- LOWER()

```
SELECT LOWER(last_name) FROM employees;
```

*** Converts characters to lowercase.*

- UPPER()

```
SELECT UPPER(last_name) FROM employees;
```

*** Converts characters to uppercase.*

- INITCAP()

```
SELECT INITCAP(last_name) FROM employees;
```

*** Capitalizes the first letter of each word.*

- CONCAT()

```
SELECT CONCAT(first_name, last_name) FROM employees;
```

```
SELECT CONCAT('Hello', 'World') FROM dual;
```

*** Joins two strings.*

With a custom string:

```
SELECT CONCAT(first_name, 'And', last_name) FROM employees;
```

- **Nested:**

```
SELECT LOWER(CONCAT(first_name, last_name)) FROM employees;
```

◆ Substring & Length Functions

- **SUBSTR()**

```
SELECT SUBSTR('HelloWorld', 1, 5) FROM dual; -- 'Hello'
```

```
SELECT SUBSTR('HelloWorld', 6) FROM dual; -- 'World'
```

*** Extracts substring starting at position with optional length.*

- **LENGTH()**

```
SELECT LENGTH('HAHAHA') FROM dual;
```

```
SELECT LENGTH(last_name) FROM employees;
```

*** Returns number of characters.*

◆ Position Function

- **INSTR()**

```
SELECT INSTR('HelloWorld', 'W') FROM dual; -- 6
```

```
SELECT last_name, INSTR(last_name, 'a') FROM employees;
```

*** Finds the position of a substring.*

Q: Find the position of letter b in the first name of employees ?

```
Select First_name,instr(First_name, 'b') from employees;
```

◆ Padding Functions

LPAD(): Pads the left side of a string.

```
SELECT LPAD('HelloWorld', 15, '-') FROM dual;
```

*** result must be total 15 letters so if it's less will add the dashes.*

Output: -----HelloWolrd

```
SELECT LPAD(last_name, 10, '*') FROM employees;
```

*** Total 10 digits or chars.*

Output: *****abel

RPAD(): Pads the Right side of a char string

```
SELECT RPAD('Hind', 8, '3') FROM dual;
```

Output : Hind3333

```
SELECT RPAD('rahma', 20, 'Oman') FROM dual;
```

Output: rahmaOmanOmanOman

◆ Trimming Characters

TRIM() : Remove all specific chars from beginning , end , or both

```
SELECT TRIM(LEADING 'a' FROM 'abcba') FROM dual;
```

-- 'bcba'

```
SELECT TRIM(TRAILING 'a' FROM 'abcba') FROM dual;
```

-- 'abcb'

```
SELECT TRIM(BOTH 'a' FROM 'abcba') FROM dual;
```

-- 'bcb'

*** Removes specific characters from start, end, or both sides.*

◆ Replacing Text

REPLACE(): Replaces characters in a string.

```
SELECT REPLACE('Jack and Joe', 'J', 'Bl') FROM dual;
```

-- 'Black and Bloe'

```
SELECT REPLACE('HOOOWORLD', 'O') FROM dual;
```

^ will delete O (HWRLD) “Can't replace O with O “

-- 'HWRLD'

■ User Input, Rounding, Date Functions in SQL

```
SELECT first_name, last_name, salary, department_id
```

```
FROM employees
```

WHERE department_id = :enter_dept_id;

**** Prompts the user to input a department ID before executing the query.**

```
SELECT tid, tname, salary
```

```
FROM teacherhct
```

WHERE tid = :enter_teacher_id;

**** User enters the Teacher ID during execution.**

Numeric Functions

- ◆ **ROUND(column, decimal_places)**

| Example | Output | Note |
|--------------------|--------|--------------------------|
| ROUND(1.95, 0) | 2 | Round to nearest integer |
| ROUND(1.94, 1) | 1.9 | Round to 1 decimal |
| ROUND(1.94, 2) | 1.94 | Already precise |
| ROUND(125, -1) | 130 | Round to nearest tens |
| ROUND(123.456, -1) | 120 | Round left of decimal |

- ◆ **TRUNC(column, decimal_places)**

| Example | Output | Note |
|-------------------|--------|-------------------------|
| TRUNC(45.926, 2) | 45.92 | Keep 2 decimals |
| TRUNC(45.926, 1) | 45.9 | Keep 1 decimal |
| TRUNC(45.926, 0) | 45 | Truncate to integer |
| TRUNC(45.926, -1) | 40 | Truncate to nearest 10 |
| TRUNC(45.926, -2) | 0 | Truncate to nearest 100 |

- ◆ **MOD(n, d) – Modulus Operator**

SELECT MOD(6, 2) FROM dual; -- 0

SELECT MOD(5, 2) FROM dual; -- 1

*** Returns the remainder after division.*

System Date and Calculations

```
SELECT SYSDATE FROM dual;
```

*** Shows current system date.*

```
SELECT last_name, hire_date + 10 FROM employees;
```

*** Adds 10 days to each hire date.*

```
SELECT last_name, (SYSDATE - hire_date) / 7 AS weeks
```

```
FROM employees;
```

*** Returns time since hire in weeks.*

```
SELECT last_name, (SYSDATE - hire_date) / 365 AS years
```

```
FROM employees;
```

*** Returns time since hire in years.*

Date Functions

- **MONTHS_BETWEEN(date1, date2)**

*** Number of months between two dates.*

```
SELECT last_name, hire_date, MONTHS_BETWEEN(SYSDATE, hire_date)
```

```
FROM employees;
```

```
SELECT MONTHS_BETWEEN('12-OCT-2021', '12-OCT-2020') FROM dual; -- 12
```

- **Filter example:**

```
SELECT last_name, hire_date
```

```
FROM employees  
WHERE MONTHS_BETWEEN(SYSDATE, hire_date) > 24;
```

- ADD_MONTHS(date, number_of_months)

```
SELECT ADD_MONTHS(SYSDATE, 1) FROM dual;
```

-- Next month from current date

```
SELECT ADD_MONTHS(SYSDATE, 12) AS next_year FROM dual;
```

- NEXT_DAY(date, 'DayName')

```
SELECT NEXT_DAY(SYSDATE, 'Saturday') AS "Next Saturday" FROM dual;
```

```
SELECT NEXT_DAY(SYSDATE, 'Monday') AS "Next Monday" FROM dual;
```

*All week days are Case sensitive.

- ◆ LAST_DAY(date)

```
SELECT LAST_DAY(SYSDATE) AS "Last Day in Month" FROM dual;
```

```
SELECT LAST_DAY(ADD_MONTHS(SYSDATE, 1)) AS "Last Day in Next Month" FROM dual;
```

Date Rounding

- ROUND(date, 'Month'):

Month=30 days

*** If day < 16 → rounds down, else → next month*

```
SELECT ROUND(DATE '2020-02-12', 'Month') FROM dual;
```

-- 01-FEB-2020

```
SELECT ROUND(DATE '2020-02-17', 'Month') FROM dual;
```

-- 01-MAR-2020

- ROUND(date, 'Year')

Year=12 months

*** If month <6 → rounds to Jan this year, else → Jan next year*

```
SELECT ROUND(DATE '2020-02-12', 'Year') FROM dual;
```

-- 01-JAN-2020

```
SELECT ROUND(DATE '2020-08-17', 'Year') FROM dual;
```

-- 01-JAN-2021

CH6.5: Date Truncation & Conversion Functions

TRUNC(date, 'format')

- Within the Month

Select hire_date, trunc(hire_date, 'Month') from employees

Where department_id=50;

* Example Output:

16-NOV-1999 → 01-NOV-1999

17-OCT-1995 → 01-OCT-1995

- Within the Year

```
SELECT hire_date, TRUNC(hire_date, 'Year')
```

FROM employees

WHERE department_id = 50;

*** Truncates the date to the first day of the same year.*

* Example Output:

16-NOV-1999 → 01-JAN-1999

17-OCT-1995 → 01-JAN-1995

■ Date → Character Conversion (TO_CHAR)

Format:

TO_CHAR(date_column, 'format_model')

Select **To_Char(Hire_date,'Month dd, yyyy')** from employees;

^ Output (June 07, 1994)

Common Date Format Models

| Format | Output Example | Description |
|--------|----------------|-------------------------|
| Month | June | Full month name |
| MON | JUN | 3-letter abbreviation |
| DD | 07 | Day of month (2 digits) |
| D | 3 | Day of week (1–7) |
| DDD | 322 | Day of the year |
| YYYY | 2024 | Year |
| HH24 | 14 | Hour (24-hour) |
| MI | 53 | Minutes |
| SS | 11 | Seconds |
| AM/PM | AM | Morning / Evening |

| Format | Output Example | Description |
|--------|--------------------|--------------------|
| SP | Seven | Spells out numbers |
| TH | 7th | Ordinal form |
| FM | Removes leading 0s | (e.g., 07 → 7) |

◆ Examples

Select `To_char (hire_date , 'month, ddth,yyyy')` from employees;

Select `To_char (sysdate , 'month , ddspth , yyyy')` from dual;

Select `To_char (hire_date , 'MON , ddth , year ')` from employees;

Select `To_char (hire_date , 'mm , ddth , year ')` from employees;

Select `To_char (hire_date , 'fmMMfm ddth , year ')` from employees;

****Note:**

DDD --> Day of the week

DD --> Day of the month

D --> day of the week

- Detailed Date + Week Formatting

`SELECT TO_CHAR(hire_date, 'DY fmMMfm ddth, Year')` FROM employees;

`SELECT TO_CHAR(hire_date, 'DDD fmMMfm ddth, Year')` FROM employees;

`SELECT TO_CHAR(SYSDATE, 'DDD fmMMfm ddth, yyyy')` FROM dual;

-- Output: 322 11 17th, 2020

`SELECT TO_CHAR(SYSDATE, 'DD fmMMfm ddth, yyyy')` FROM dual;

-- Output: 17 11 17th, 2020

-
- Time Display with AM/PM and 24-Hour Clock

```
SELECT TO_CHAR(SYSDATE, 'HH24-MI-SS PM - mm, dd, yyyy') FROM dual;
```

-- Output: 10-49-11 AM - 11, 17, 2020

```
SELECT TO_CHAR(SYSDATE, 'HH-MI-SS - mm, dd, yyyy') FROM dual;
```

-- Output: 10-50-44 - 11, 17, 2020

*** Even if you use AM/PM, the system time will reflect actual system clock.*

- TO_CHAR(date, 'format') with Custom Text

```
SELECT TO_CHAR(SYSDATE, 'DD "OF" MONTH') FROM dual;
```

-- Output: 17 OF NOVEMBER

*** You can include any static characters inside double quotes within a format string.*

■ Number → Character Formatting

- Format: '9'

```
SELECT TO_CHAR(1234, '99999') FROM dual;
```

-- Output: 1234

*** 9 is a digit placeholder that doesn't add leading zeros.*

- Format: '0'

```
SELECT TO_CHAR(1234, '099999') FROM dual;
```

-- Output: 001234

*** 0 adds leading zeros — 1 zero adds 2 digits, 2 zeros add 4, and so on.*

- Format: '\$'

```
SELECT TO_CHAR(1234, '$099,999') FROM dual;
```

-- Output: \$001,234

*** Includes a dollar sign and comma formatting.*

- Format: 'L' (Local currency)

```
SELECT TO_CHAR(1234, 'L99,99') FROM dual;
```

-- Output: \$12,34 (depending on NLS settings)

```
SELECT TO_CHAR(1234.333, 'L99,99') FROM dual;
```

*** Might fail or round depending on decimals & width.*

- Format: 'MI' (Negative after number)

```
SELECT TO_CHAR(-1234, '999,9MI') FROM dual;
```

-- Output: 123,4-

- Format: 'PR' (Negative in angle brackets)

```
SELECT TO_CHAR(-1234, '99999PR') FROM dual;
```

-- Output: <1234>

- Format: 'V' (Virtual decimal)

```
SELECT TO_CHAR(4.5, '9999999V99') FROM dual;
```

-- Output: 000000450

*** Moves the decimal: 4.5 → 450*

- Format: 'B' (Blank for zero)

```
SELECT TO_CHAR(040.00, 'B999999999') FROM dual;
```

-- Output: 40

*** Suppresses leading zeros.*

■ Character → Date Conversion

- TO_DATE('string', 'format')

```
SELECT TO_DATE('April 07, 2020', 'Month DD, YYYY') FROM dual;
```

```
SELECT TO_DATE('April 07, 2020', 'Mon DD, YYYY') FROM dual;
```

*** Both Month and Mon are accepted.*

- Using Fx Modifier (Format-Exact)

```
SELECT TO_DATE('Apr 07, 2020', 'FxMon DD, YYYY') FROM dual;
```

```
SELECT TO_DATE('April 07, 2020', 'FxMonth DD, YYYY') FROM dual;
```

*** Ensures strict matching (Apr ≠ April in FxMon)*

- RR vs YY in Year Format

```
SELECT TO_DATE('June 19, 04', 'Month DD, YY') FROM dual;
```

-- Output: 19-JUN-2004

```
SELECT TO_DATE('27-OCT-95', 'DD-Mon-YY') FROM dual;
```

-- Output: 27-OCT-2095

(It gives us 2095 'In this case' we'll use RR instead of YY)

```
SELECT TO_DATE('27-OCT-95', 'DD-Mon-RR') FROM dual;
```

-- Output: 27-OCT-1995

| Current Year | Specified Date | RR Format | VV Format |
|--------------|----------------|-----------|-----------|
| 1995 | 27-Oct-95 | 1995 | 1995 |
| 1995 | 27-Oct-17 | 2017 | 1917 |
| 2017 | 27-Oct-17 | 2017 | 2017 |
| 2017 | 27-Oct-95 | 1995 | 2095 |

CH7: NULL Handling & Conditional Value Functions

- **NVL(expr1, expr2)**

If expr1 is NULL, returns expr2. Otherwise, returns expr1.

Select **nvl** (**Salary**, **0**) from employees;

--> Null values will be displayed as 0's

Select **nvl(Date_of_independents , 'No Date')** from wf_country;

--> All null Dates will be Displayed as No date .

Other examples:

Select last_name,**NVL (Commision_pct,0)*250 AS "commission"**

from employees Where department_id IN(80,90);

-
- **NVL2(expr1, expr2, expr3)**

Logic:

- If expr1 is not **NULL** → returns expr2
- If expr1 is **NULL** → returns expr3

SELECT last_name, salary,

NVL2(commision_pct, salary + 5, salary) AS income FROM employees;

→ Adds 5 to salary only if commission exists. Else salary will be the same.

- **NULLIF(expr1, expr2)**

* *Returns NULL if both expressions are equal; otherwise returns expr1.*

```
SELECT first_name, LENGTH(first_name) AS "Length FN",
last_name, LENGTH(last_name) AS "Length LN",
NULLIF(LENGTH(first_name), LENGTH(last_name)) AS "Compare Them"
FROM employees;
```

Test example:

```
SELECT NULLIF(LENGTH('Hind'), LENGTH('Hind222')) FROM dual;
```

-- Output: NULL (lengths are not equal)

- **COALESCE FUNCTION :**

* *Returns the first non-NULL expression in the list.*

```
SELECT last_name,
COALESCE(commission_pct, salary, 10) AS "Comm"
FROM employees
ORDER BY commission_pct;
```

Select COALESCE(commission_pct, First_Name, Last_Name)

From employees;

-- Returns the first non-null value among the three columns.

| FUNCTION | DESCRIPTION |
|----------|--|
| NVL | Replace NULL with a default value |
| NVL2 | Choose between two values depending on whether input is null |
| NULLIF | Compare two values; return NULL if equal |
| COALESCE | Return the first non-null value from a list |

■ Conditional Expressions – CASE Statement:

CASE **expr**

WHEN **comparison_expr1** THEN **return_expr1**

WHEN **comparison_expr2** THEN **return_expr2**

...

ELSE **else_expr**

END

Example:

```
SELECT last_name,department_id,
CASE department_id WHEN 90 THEN 'Management'
                  WHEN 80 THEN 'Sales'
                  WHEN 60 THEN 'It'
                  ELSE 'Other dept.'
END AS "Department" FROM employees;
```

CH7: Joins in SQL

1. NATURAL JOIN

- * Automatically joins tables based on columns with the same name and data type
- * No need to specify the join condition manually.

```
SELECT first_name, last_name, job_title  
FROM employees  
NATURAL JOIN jobs  
WHERE department_id > 80;
```

Another example:

```
SELECT department_name, city  
FROM departments  
NATURAL JOIN locations;
```

2. CROSS JOIN

Join each row of a table to every row in the other table.

- * First Row of table1 with all the rows from table2
- * Second Row of table1 with all the rows from table2, And so on ..

```
SELECT last_name, department_name  
FROM employee  
CROSS JOIN department;
```

```
SELECT tname, person_id  
FROM teacherhct  
CROSS JOIN personal;
```

```
SELECT first_name, job_title  
FROM employees  
CROSS JOIN jobs;
```

→ If table A has 4 rows and table B has 3 → result = $4 \times 3 = 12$ rows.

3. JOIN ... USING(column_name)

Explicitly joins using a common column that has the same name in both tables.

```
SELECT first_name, last_name, department_id, department_name  
FROM employees  
JOIN departments USING (department_id);
```

- With condition:

```
SELECT first_name, last_name, department_id, department_name  
FROM employees  
JOIN departments USING (department_id)  
WHERE department_id = 10;
```

4. Aliases + JOIN / WHERE

```
SELECT e.first_name, e.last_name, j.job_id  
FROM employees e  
JOIN jobs j USING (job_id);
```

- Using traditional JOIN with WHERE:

```
SELECT last_name, e.job_id, job_title  
FROM employees e , jobs j  
WHERE e.job_id = j.job_id AND department_id = 80;
```

Another format:

```
SELECT last_name, emp.job_id, job_title AS "ddd"  
FROM employees emp, jobs j  
WHERE emp.job_id = j.job_id AND department_id = 80;
```

CH7.1: ON Clause in SQL Joins

- When to Use ON

- Use when **column names differ** between the tables.
- Use when the join needs **non-equality** conditions like **<, >, BETWEEN**.
- Unlike USING, ON allows **flexibility** in join logic.
- You can still apply **WHERE** conditions afterward.

- Syntax:

```
SELECT column_list  
FROM table1 alias1  
JOIN table2 alias2  
ON (alias1.column = alias2.column);
```

- Example from the slide:

```
SELECT last_name, job_title  
FROM employees e  
JOIN jobs j  
ON (e.job_id = j.job_id);
```

*** Required when job_id in employees ≠ job_id in jobs by name or if you need complex conditions.*

- ◆ Q1. Locations + Countries Join

- Show capitalized city, and replace 'a' with '\$' in country_name
Condition:
 - location_id > 500
 - city ends with 'e'

Answer:

```
SELECT INITCAP(city), REPLACE(country_name, 'a', '$')  
FROM locations  
NATURAL JOIN countries  
WHERE location_id > 500 AND city LIKE '%e';
```

Q2. Employee Name Format with @

- Display employee_id, and first_name with e replaced by @
Label the name column as "EMPLOYEE NAME"
Exclude IDs 103 and 100

Answer:

```
SELECT employee_id,  
       REPLACE(first_name, 'e', '@') AS "EMPLOYEE NAME"  
  FROM employees  
 WHERE employee_id NOT IN (103, 100);
```

Q3. Email + Password Generation (Advanced)

Use Employees table to generate:

- Full name (First + Last)
- Email: lowercase first_name + @hct.edu.om
- Password: last 3 letters of first_name, starting with a capital letter

Answer:

```
SELECT first_name,  
       last_name,  
       LOWER(first_name) || '@hct.edu.om' AS "Email Address",  
       INITCAP(SUBSTR(first_name, -3)) AS "Password"  
  FROM employees;
```

***Explanation:*

LOWER(first_name) → lowercase email prefix

SUBSTR(first_name, -3) → last 3 characters

INITCAP(...) → capitalize first letter of password

CH7.2: Group Functions in SQL

- SUM

```
SELECT SUM(salary)  
FROM employees;  
* Adds up all salary values.
```

- AVG (Average)

```
SELECT AVG(salary)  
FROM employees  
WHERE department_id = 10;  
* Average salary in department 10.
```

```
SELECT ROUND(AVG(salary), 2)  
FROM employees  
WHERE department_id = 90;  
* Rounded average to 2 decimal places.
```

◆ COUNT

```
SELECT COUNT(*) FROM employees;  
-- Counts all rows including NULLs
```

```
SELECT COUNT(hire_date) FROM employees;  
-- Counts only rows with non-null hire_date  
SELECT COUNT(job_id) FROM employees;
```

```
SELECT COUNT(commission_pct) FROM employees;
```

* Ignores NULL values.

- VARIANCE & STDDEV

```
SELECT VARIANCE(salary) FROM employees;
```

```
SELECT STDDEV(salary) FROM employees;
```

* Measures how spread out salary values are — standard deviation & variance.

* Helps understand deviation from the average.

- MAX, MIN (Multiple Group Functions)

```
SELECT MAX(salary), MIN(salary), MIN(employee_id)
```

```
FROM employees
```

```
WHERE department_id = 6;
```

**** we cant use do :**

```
SELECT max(salary), min(salary), min (employee_id), employee_id
```

```
From employees
```

```
Where department_id = 6
```

X You cannot mix group functions with individual row columns unless using GROUP BY.

COUNT with DISTINCT

- COUNT(*)

** Counts all rows — even with nulls.*

```
SELECT COUNT(*) FROM employees;
```

-- Output: total number of rows

- COUNT(column)

** Counts only non-null values.*

```
SELECT COUNT(commission_pct) FROM employees;
```

-- Output: 4 (if 6 are null)

- COUNT(DISTINCT column)

** Counts non-duplicate, non-null values.*

```
SELECT COUNT(DISTINCT job_id)
```

```
FROM employees;
```

- DISTINCT Keyword

** Removes duplicate values from result set.*

```
SELECT DISTINCT job_id
```

```
FROM employees;
```

```
SELECT DISTINCT hire_date
```

```
FROM employees
```

```
ORDER BY hire_date;
```

CH8: GROUP BY & HAVING

- GROUP BY Clause

Groups rows based on one or more columns, so you can apply aggregate functions (MAX, COUNT) to each group.

Examples

```
SELECT MAX(salary)  
FROM employees  
GROUP BY department_id;
```

Returns the highest salary per department.

```
SELECT MAX(salary), department_id  
FROM employees  
GROUP BY department_id;
```

Now we see both the department ID and its max salary.

-
- Another grouping example:

```
SELECT MAX(salary)  
FROM employees  
GROUP BY manager_id;
```

Try this to see max salary per manager ID.

Note:

- You can't use WHERE to filter grouped values, that's what HAVING is for.
 - WHERE filters rows before grouping, while HAVING filters groups after aggregation.
-

- HAVING Clause

| Clause | Filters What? | Used With |
|--------|-----------------------------------|--------------------|
| WHERE | Individual rows (before grouping) | All queries |
| HAVING | Grouped results (after GROUP BY) | Only with GROUP BY |

Examples

```
SELECT department_id, job_id, COUNT(*)
```

```
FROM employees
```

```
GROUP BY department_id, job_id
```

```
HAVING COUNT(*) > 2;
```

** Shows departments & jobs that appear more than twice.*

```
SELECT department_id, job_id, COUNT(*)
```

```
FROM employees
```

```
GROUP BY department_id, job_id
```

```
HAVING department_id > 50;
```

** Filters grouped rows where department_id > 50.*

- More Examples:

```
SELECT department_id, MAX(salary)  
FROM employees  
GROUP BY department_id  
HAVING COUNT(*) > 1  
ORDER BY department_id;
```

```
SELECT department_id, MAX(salary)  
FROM employees  
GROUP BY department_id  
HAVING MAX(salary) = 5800  
ORDER BY department_id;
```

* This shows only departments where the highest salary = 5800.

CH9: Subqueries and Multi-Row Comparison Operators

What is a Subquery?

A subquery is a query nested inside another query's WHERE, FROM, or SELECT.

Example: Simple Subquery with >

```
SELECT first_name, last_name, hire_date FROM employees  
WHERE hire_date > (  
    SELECT hire_date  
    FROM employees  
    WHERE last_name = 'Vargas' );
```

* Returns employees hired after Vargas.

1. ANY Operator

* Compares a value to any value returned by the subquery.

* Return rows if the condition is true for at least one value.

```
SELECT salary  
FROM employees  
WHERE salary > ANY (  
    SELECT salary  
    FROM employees  
    WHERE department_id = 20  
);
```

~ If department 20 has salaries 3000, 3500, 4000 →

It returns salaries higher than at least one of these (i.e., >3000).

2. ALL Operator

* Compares a value to all values returned by the subquery.

* Condition must be true for every value in the result set.

```
SELECT salary  
FROM employees  
WHERE salary > ALL (  
    SELECT salary  
    FROM employees  
    WHERE department_id = 20  
);
```

Only returns salaries higher than everyone in dept 20.

3. IN Operator

* Matches exact values from the subquery.

```
SELECT last_name, employee_id  
FROM employees  
WHERE employee_id IN (  
    SELECT manager_id  
    FROM employees  
);
```

* Returns employees who are managers (since their IDs appear as manager_id).

Example2:

```
SELECT first_name, last_name, hire_date  
FROM employees  
WHERE first_name IN (  
    SELECT first_name  
    FROM employees  
    WHERE first_name LIKE '%a%'  
);
```

→ Finds employees whose first names contain 'a' — since the inner query collects all such names.

CH10: SQL Views

- What is a View?

A view is a virtual table based on a SQL query.

It can:

- Show specific columns from a table
 - Apply restrictions to limit visible data
 - Be updated like a table (with some rules)
-

1. Create a View

```
CREATE VIEW view_employees AS  
SELECT employee_id, first_name, last_name, email, phone_number  
FROM employees;  
* Creates a view named view_employees with selected columns only.
```

```
SELECT * FROM view_employees;
```

2. Create or Replace View

```
CREATE OR REPLACE VIEW view_employees AS  
SELECT employee_id, first_name, last_name  
FROM employees;  
* Replaces the view if it already exists, with a new column set.
```

- Simple View

- Has One base table
- No functions, group by, joins, etc.

We *can also insert into simple views:*

```
INSERT INTO view_employees VALUES (123, 'Ali', 'Ahmed');
```

* Works only if all columns are from one base table and no constraints are violated.

- Complex View has:

- Joins
- Group by
- Aggregates (e.g., MIN, MAX)

```
CREATE VIEW view_123 AS
```

```
SELECT employee_id, job_id, job_title
```

```
FROM employees
```

```
JOIN jobs USING (job_id);
```

~ This view cannot be updated/inserted directly.

| Feature | Simple Views | Complex Views |
|--|--------------|---------------|
| Number of tables used to derive data | One | One or more |
| Can contain functions | No | Yes |
| Can contain groups of data | No | Yes |
| Can perform DML operations (INSERT, UPDATE, DELETE) through a view | Yes | Not always |

CH11.1: GRANT – Giving Privileges

1. Grant SELECT to All Users (PUBLIC)

GRANT SELECT ON view_employees TO PUBLIC;

~ Gives read access to all users.

2. To view the shared data:

SELECT * FROM **username.tablename**;

SELECT * FROM **om_a835_sql_s57.view_employees**;

3. Grant SELECT to a Specific User

GRANT SELECT ON **view_employees** TO **om_a102_sql_s40**;

~ Now only that user can read the view.

4. Grant UPDATE Privileges

GRANT UPDATE ON view_employees TO PUBLIC;

GRANT **UPDATE** ON **view_employees** TO **om_a102_sql_s40**;

~ Cannot update complex views, only simple views.

5. Grant DELETE Privileges

GRANT DELETE ON view_employees TO PUBLIC;

GRANT **DELETE** ON **view_employees** TO **om_a102_sql_s40**;

6. Grant ALL Privileges (Insert, Select, Update, Delete)

GRANT **INSERT, SELECT, UPDATE, DELETE**

ON **view_employees** TO PUBLIC;

CH11.2: REVOKE – Removing Privileges

1. Revoke SELECT

```
REVOKE SELECT ON view_employees FROM PUBLIC;
```

2. Revoke UPDATE

```
REVOKE UPDATE ON view_employees FROM PUBLIC;
```

```
REVOKE UPDATE ON view_employees FROM om_a102_sql_s40;
```

3. Revoke DELETE

```
REVOKE DELETE ON view_employees FROM PUBLIC;
```

```
REVOKE DELETE ON view_employees FROM om_a102_sql_s40;
```

4. Revoke ALL Privileges

```
REVOKE INSERT, SELECT, UPDATE, DELETE
```

```
ON view_employees FROM PUBLIC;
```

CH11.2: Transaction Control in SQL

COMMAND DESCRIPTION

| | |
|--------------|--|
| SAVEPOINT T1 | Marks a point to return to |
| ROLLBACK T1 | Undoes changes after T1 |
| COMMIT | Finalizes all changes, cannot rollback |

Example of savepoint:

At Step 3 – Update Ford to BMW → SAVEPOINT T1

UPDATE car

SET carname = 'BMW'

WHERE carno = 200;

SAVEPOINT T1;

At Step 4 – Insert Lexus → SAVEPOINT T2

INSERT INTO car VALUES(300, 'Lexus');

SAVEPOINT T2;

At Step 5 – Update Honda to GMC → SAVEPOINT T3

UPDATE car

SET carname = 'GMC'

WHERE carno = 100;

SAVEPOINT T3;

Example of ROLLBACK:

ROLLBACK TO T2;

- **Keeps** changes up to T2 (BMW and Lexus insert stay)
 - **Undoes** the update to GMC and deletion of car 200
-

- COMMIT:

COMMIT;

~ Locks in all changes made before — no rollback possible after this.